

# Clairvoyant Prefetching for Machine Learning I/O

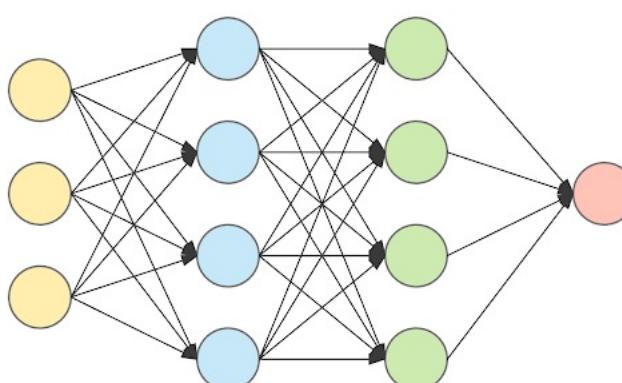
**ROMAN BÖHRINGER**



# Supervised Deep Learning



$$f(x; \theta)$$



$$f^*(x)$$

0  
1  
2  
3  
4  
5  
6  
7  
8  
9

---

**Algorithm 1:** Stochastic Gradient Descent [16].

---

**Require:** Learning rate schedule  $\varepsilon_1, \varepsilon_2, \dots$

**Require:** Initial parameter  $\theta$

$k \leftarrow 1;$

**while** stopping criterion not met **do**

Sample mini-batch of  $m$  examples from the training set  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  with corresponding targets  $\mathbf{y}^{(i)}$ ;

Compute gradient estimate:  $\hat{\mathbf{g}} \leftarrow \frac{1}{m} \nabla_{\theta} \sum_i L(f(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)})$ ;

Apply update:  $\theta \leftarrow \theta - \varepsilon_k \hat{\mathbf{g}}$ ;

$k \leftarrow k + 1$ ;

**end**

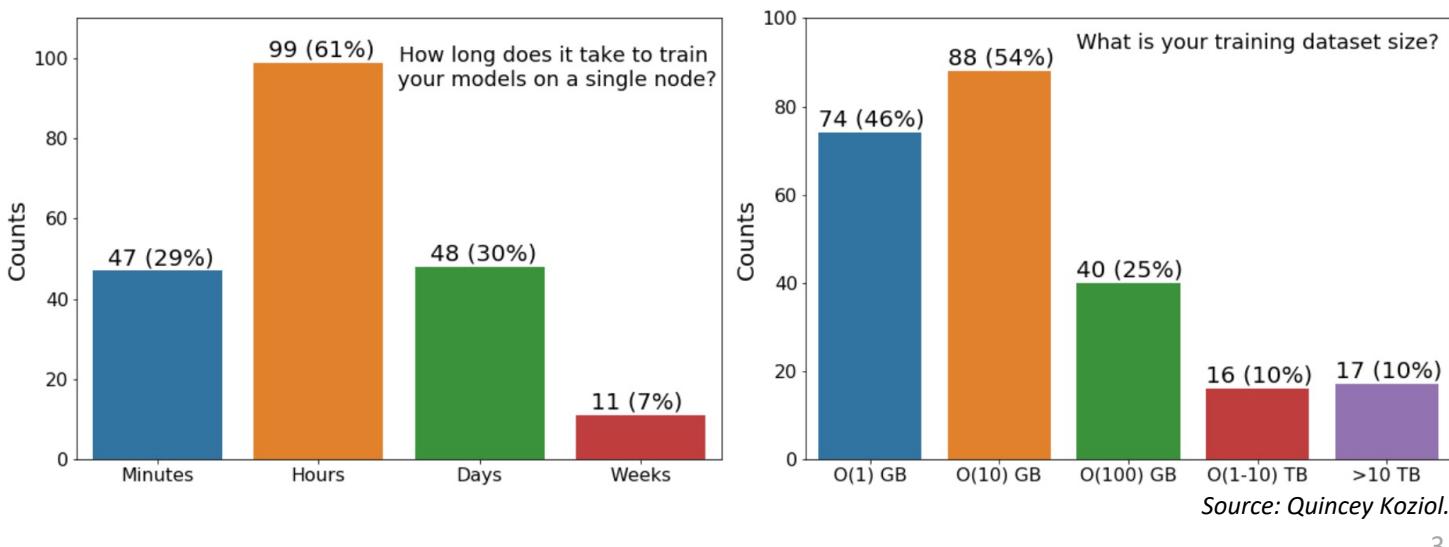
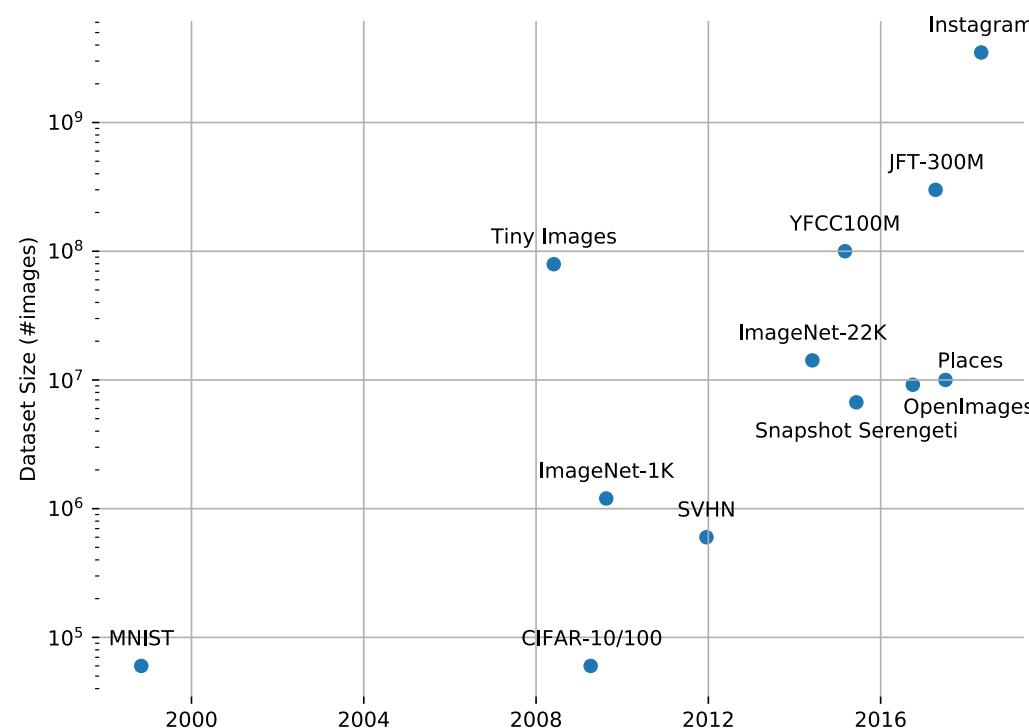
---

# Deep Learning Trends & Challenges

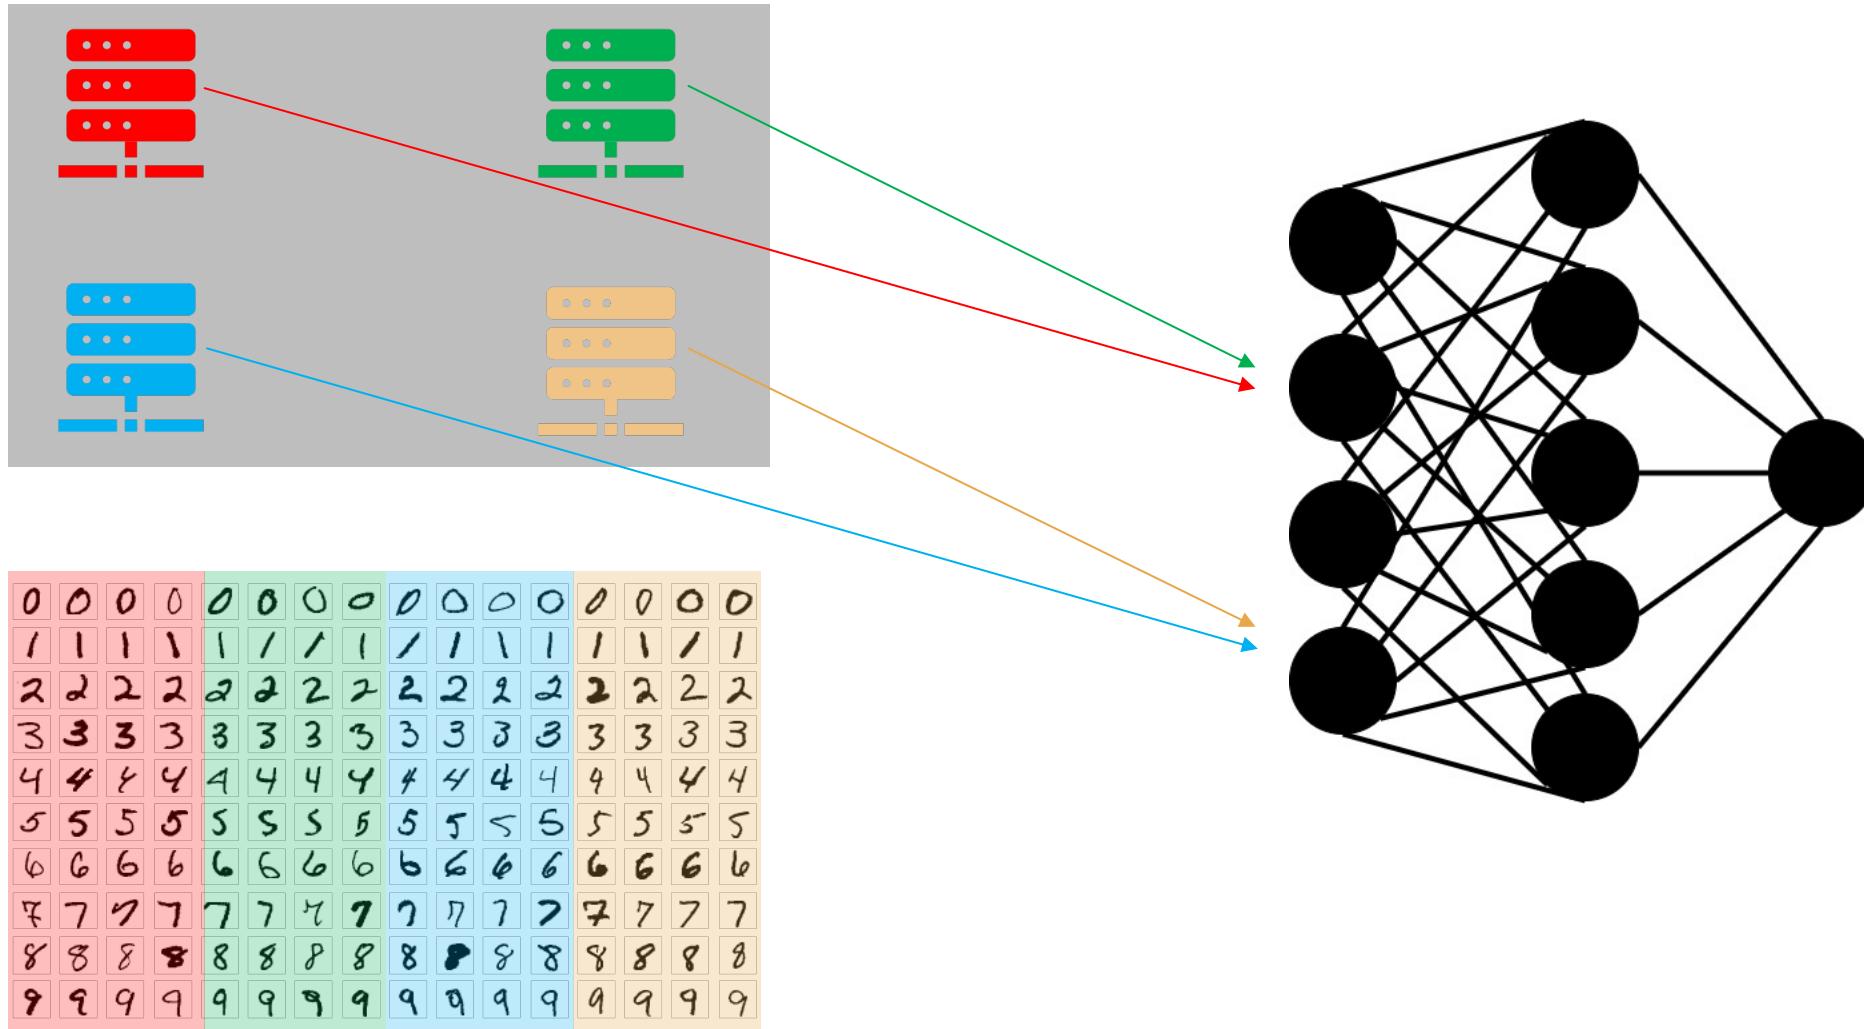
Model Name	$n_{\text{params}}$	$n_{\text{layers}}$	$d_{\text{model}}$	$n_{\text{heads}}$	$d_{\text{head}}$	Batch Size	Learning Rate
GPT-3 Small	125M	12	768	12	64	0.5M	$6.0 \times 10^{-4}$
GPT-3 Medium	350M	24	1024	16	64	0.5M	$3.0 \times 10^{-4}$
GPT-3 Large	760M	24	1536	16	96	0.5M	$2.5 \times 10^{-4}$
GPT-3 XL	1.3B	24	2048	24	128	1M	$2.0 \times 10^{-4}$
GPT-3 2.7B	2.7B	32	2560	32	80	1M	$1.6 \times 10^{-4}$
GPT-3 6.7B	6.7B	32	4096	32	128	2M	$1.2 \times 10^{-4}$
GPT-3 13B	13.0B	40	5140	40	128	2M	$1.0 \times 10^{-4}$
GPT-3 175B or "GPT-3"	175.0B	96	12288	96	128	3.2M	$0.6 \times 10^{-4}$

**Table 2.1:** Sizes, architectures, and learning hyper-parameters (batch size in tokens and learning rate) of the models which we trained. All models were trained for a total of 300 billion tokens.

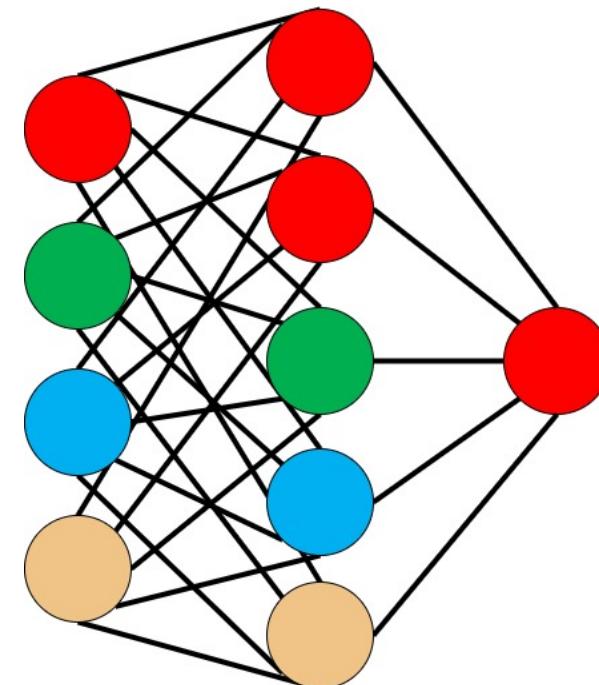
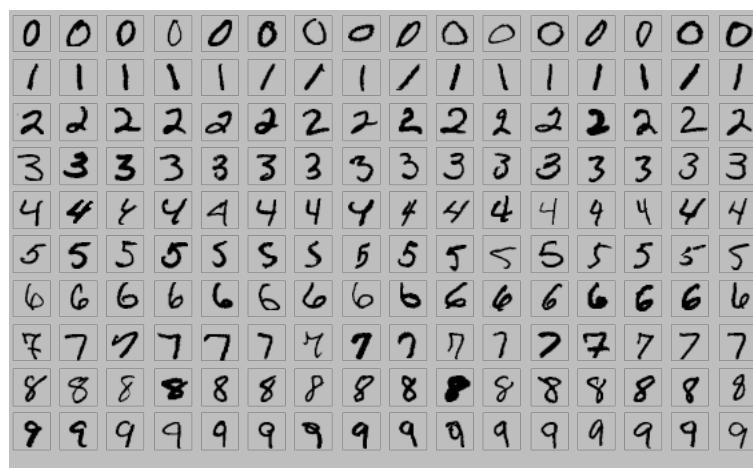
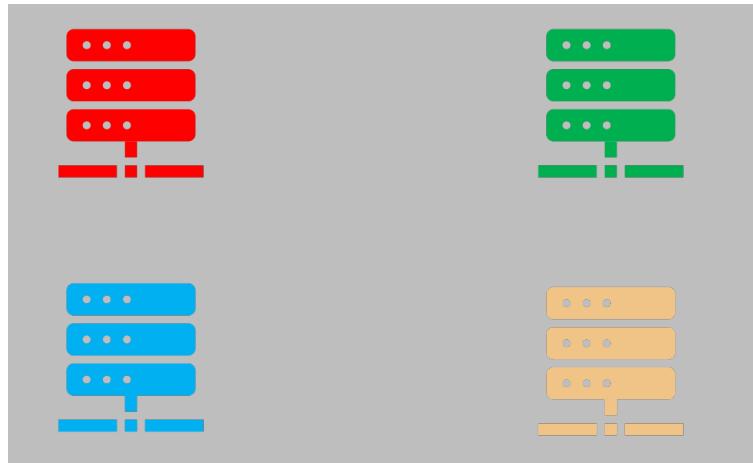
Source: Tom B. Brown et al.



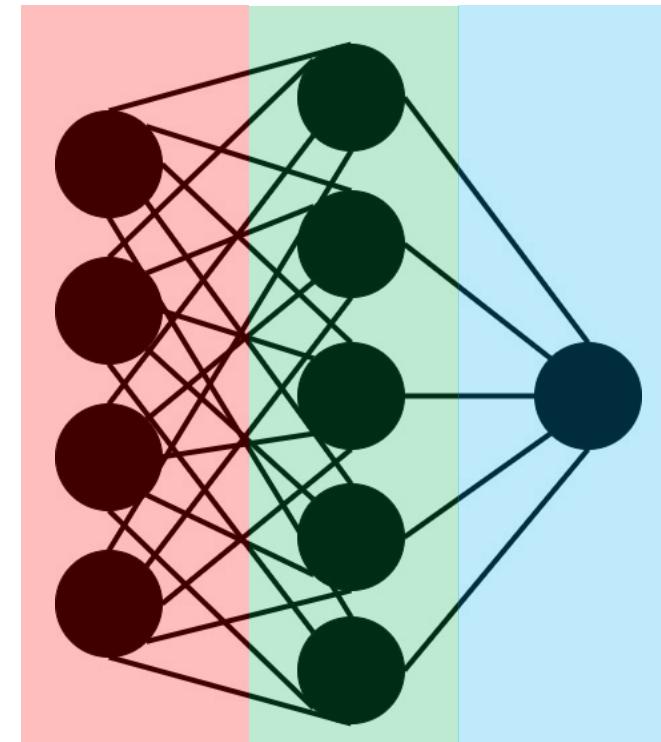
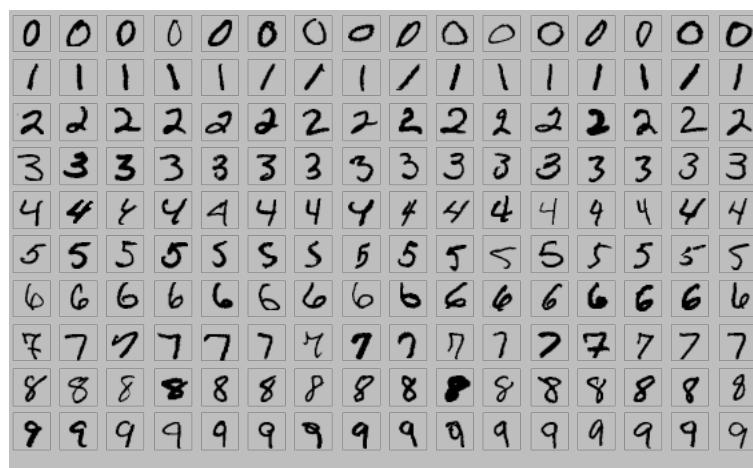
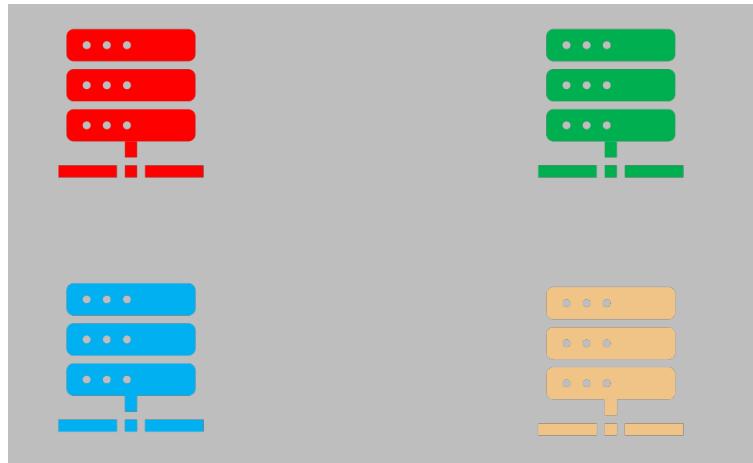
# Distributed Deep Learning: Data Parallelism



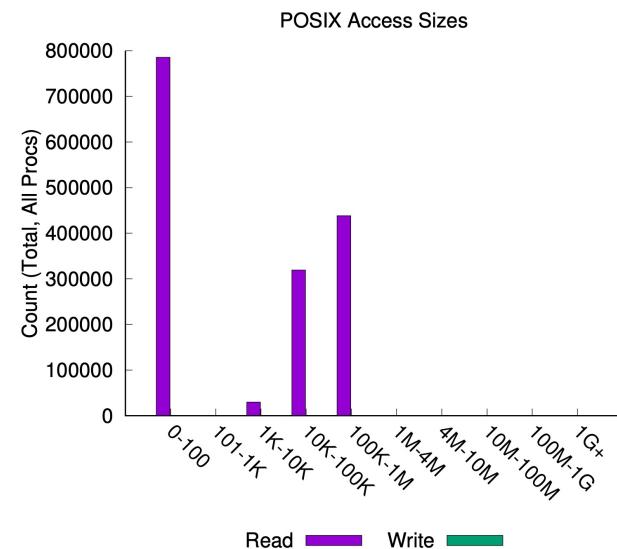
# Distributed Deep Learning: Model Parallelism



# Distributed Deep Learning: Pipelining



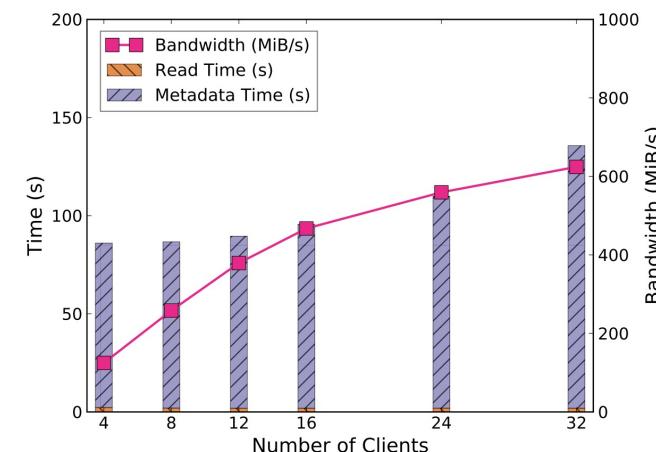
# (Distributed) Deep Learning: I/O



(b) Operation Count vs. File Sizes

Source: Fahim Chowdhury et al.

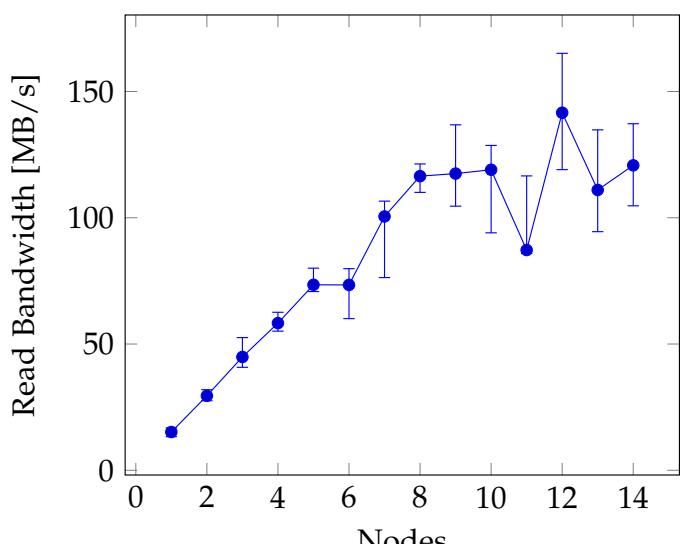
TensorFlow ImageNet Data Reader I/O Pattern.



I/O Latency and Read Bandwidth of ImageNet Data Reader Pipeline on TensorFlow.

## HPC FILE SYSTEMS FAIL FOR DEEP LEARNING AT SCALE

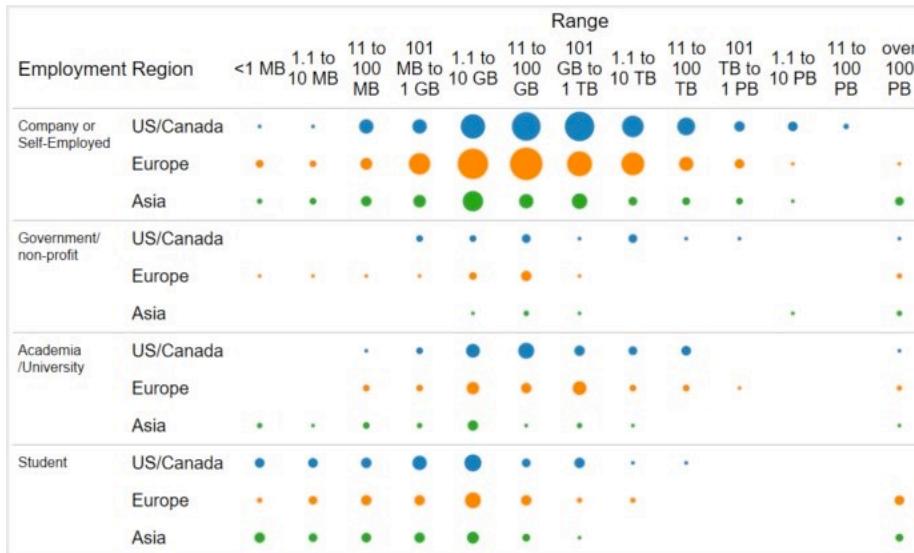
October 9, 2018 Nicole Hemsoth



IOR Benchmark Results.

# Deep Learning I/O Challenges

## 1. Dataset Scalability



Source: Gregory Piatetsky.

## 2. Node Scalability

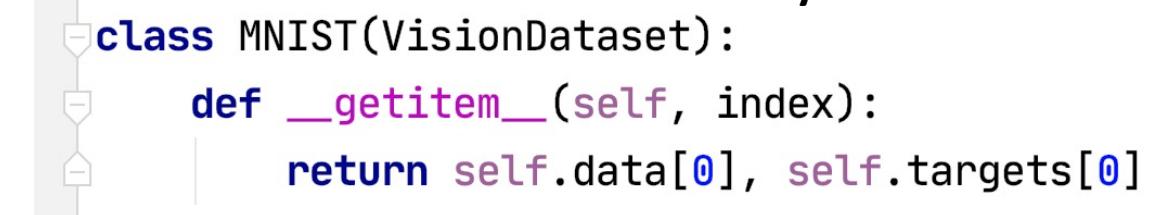
Cores	Steps per sec.	Batch sz. (Tokens)	TPU core years	Training time (days)
2048	0.72	4M	22.4	<b>4.0</b>
2048	2.15	4M	7.5	1.4
512	1.05	1M	15.5	11.0
512	3.28	1M	4.9	3.5
128	0.67	1M	6.1	17.3
128	2.16	1M	1.9	5.4
<b>2048</b>	-	4M	~235.5	~42

Source: Dmitry Lepikhin et al.

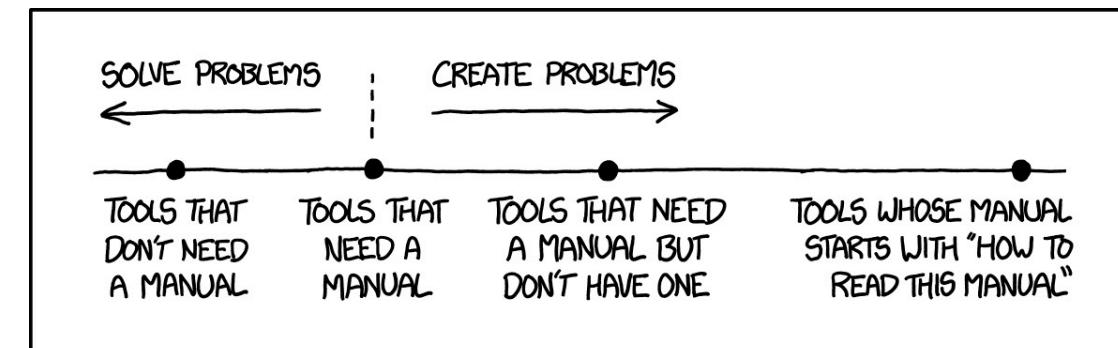
## 3. Configuration Independence

System	Node Local Storage
Piz Daint	No
SuperMUC-NG	No
Sierra	1.6TB NVMe PCIe SSD
Summit	800GB non-volatile memory

## 4. Model Accuracy



## 5. Ease of Use



Source: Randall Munroe.

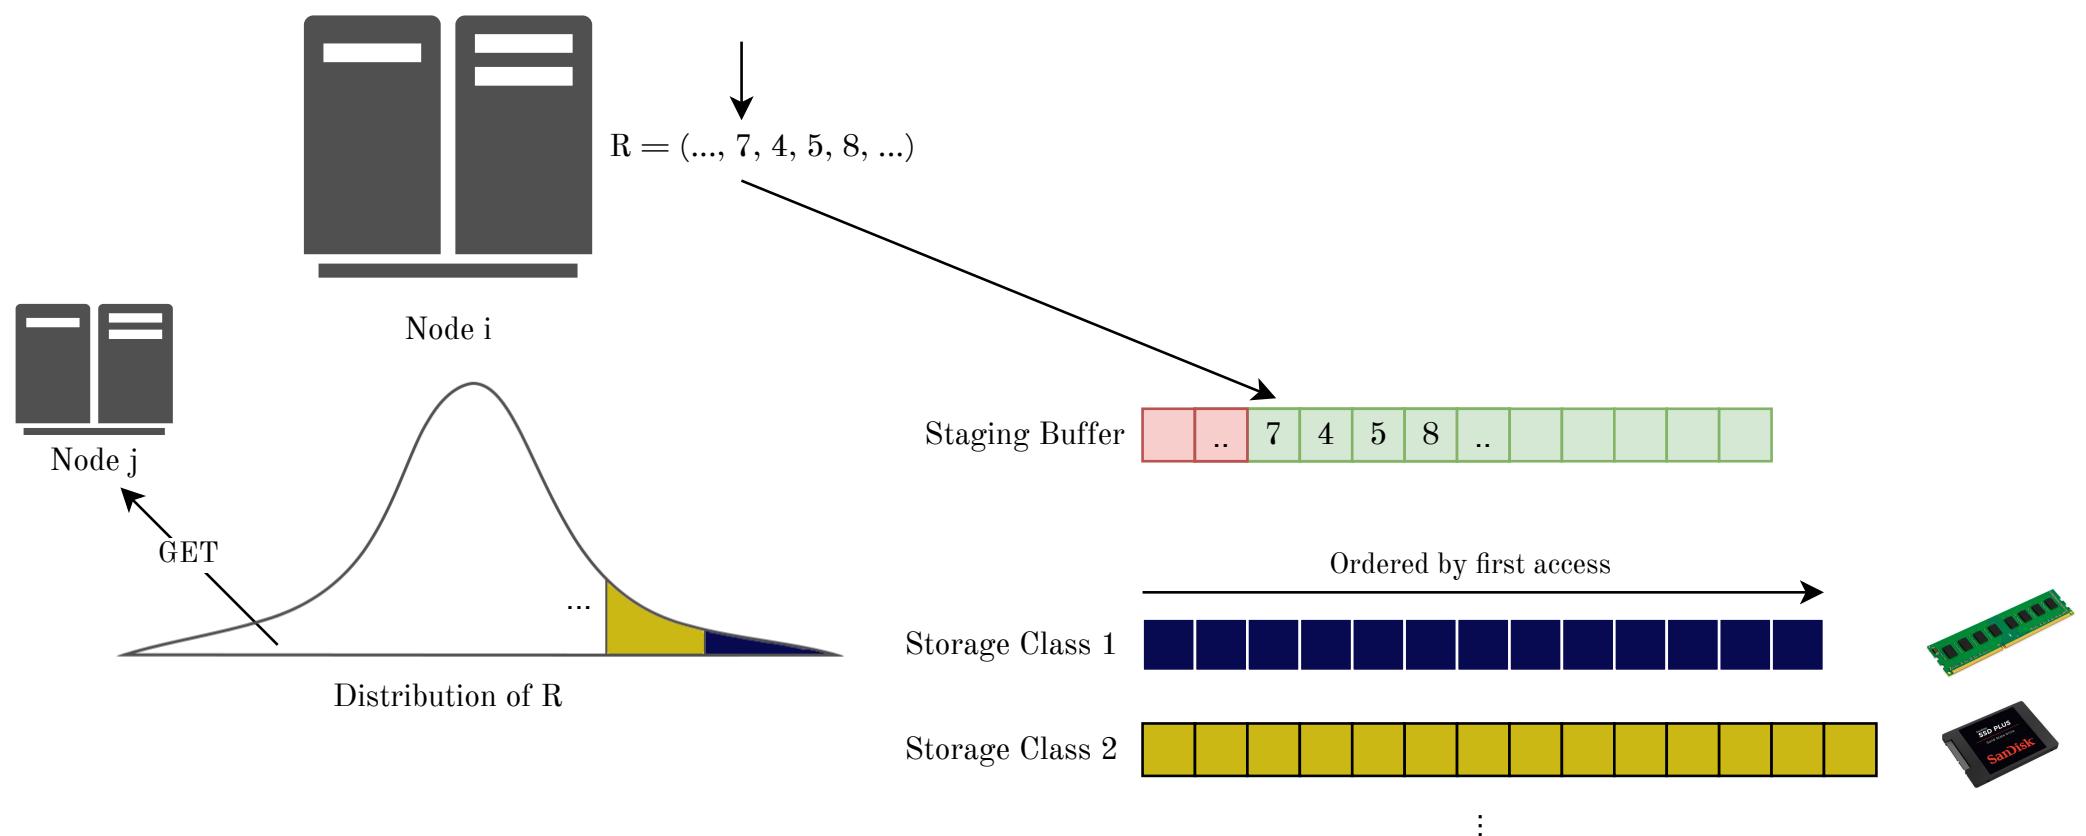
# Deep Learning I/O: Solutions

Approach	Dataset Scalability	Node Scalability	Configuration Independence	Model Accuracy	Ease of Use
tf.data with shuffle/prefetch [12]	✓	✗	✗	✗	✓
Double Buffering	✓	✗	✗	✓	✓
DeepIO [47]	✓	✓	✗	✗ <sup>a</sup>	✓
Parallel Data Staging [28]	✓	✓	✗	✗	✗
LBANN Distributed Data Store [21]	✗	✓	✗	✓	✗
Locality-Aware Data Loading [44]	✓	✓	✗	✓ <sup>b</sup>	? <sup>c</sup>

<sup>a</sup>In entropy-aware opportunistic mode

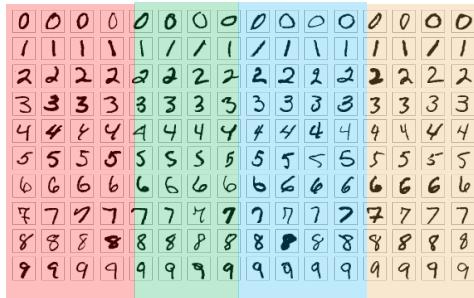
<sup>b</sup>Result can be different when using node-local batch normalization

<sup>c</sup>PyTorch based prototype, general software package is planned according to the paper

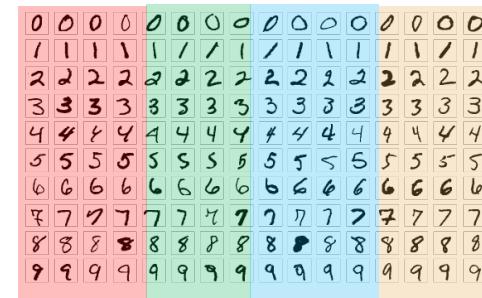
 HDMLP

# Data Parallelism: Access Frequency

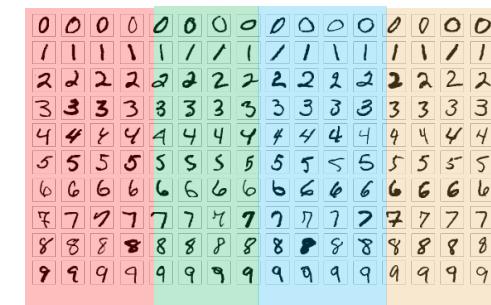
## Epoch 1



## Epoch 2

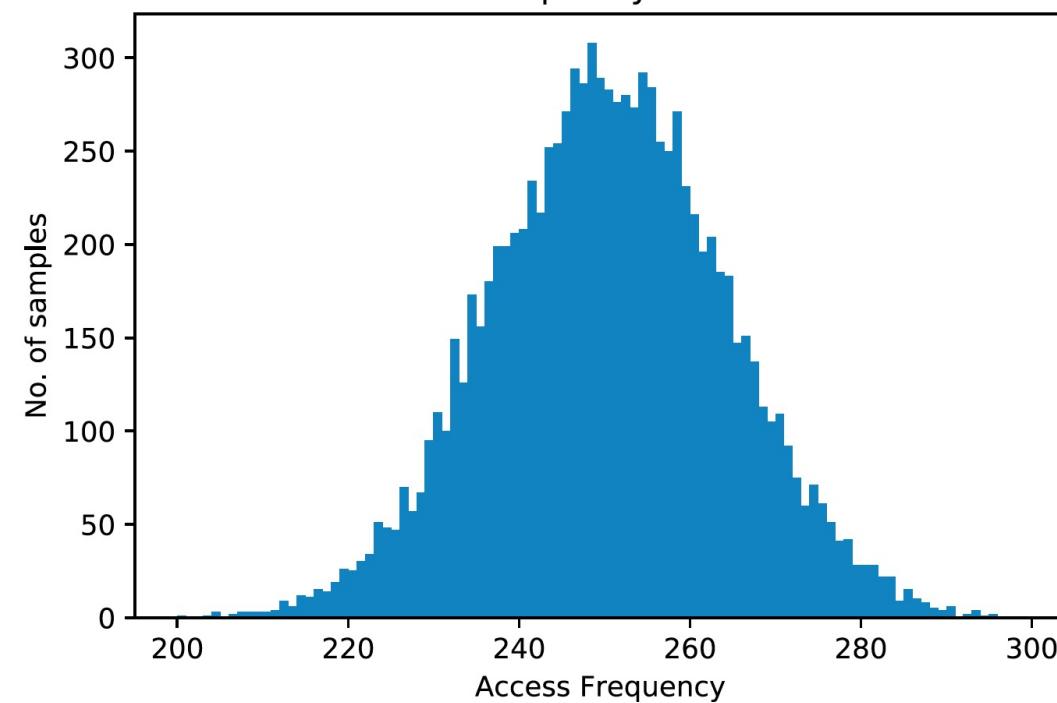


## Epoch E



•

## Access Frequency distribution

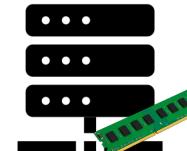


# Challenges 1 – 4: Performance Model

Variable	Unit	Definition
$N$		Number of nodes
$c$	MB/s	Compute throughput
$b_c$	MB/s	Network bandwidth for inter-node communication
$b_{fs}$	MB/s	Network bandwidth for communication with the PFS
$p_j$		Number of threads for prefetching to storage class $j$
$d_j$	MB	Capacity of storage class $j$
$D$	MB	Total local storage of a node
$r_j(p)$	MB/s	Random aggregate read throughput of storage class $j$ (with $p$ reader threads)
$w_j(p)$	MB/s	Random aggregate write throughput of storage class $j$ (with $p$ writer threads)
$t(\gamma)$	MB/s	Random aggregate read throughput (with $\gamma$ clients) of the PFS
$\beta$	MB/s	Preprocessing rate
$F$		Number of files
$s_k$	MB	Size of sample $k$
$S$	MB	Size of the whole dataset
$B$		Batch size
$E$		Number of epochs
$T$		Number of iterations
$B^h$		Subset of samples that are processed at iteration $h$
$B^{h,i}$		Subset of samples that are processed by node $i$ at iteration $h$
$R$		Access pattern of a node

# Performance Model: Scenarios

## 1. Whole dataset fits into memory of a node



## 2. Dataset fits into aggregated storage of a node



### **3. Dataset is smaller than aggregated storage of all nodes**



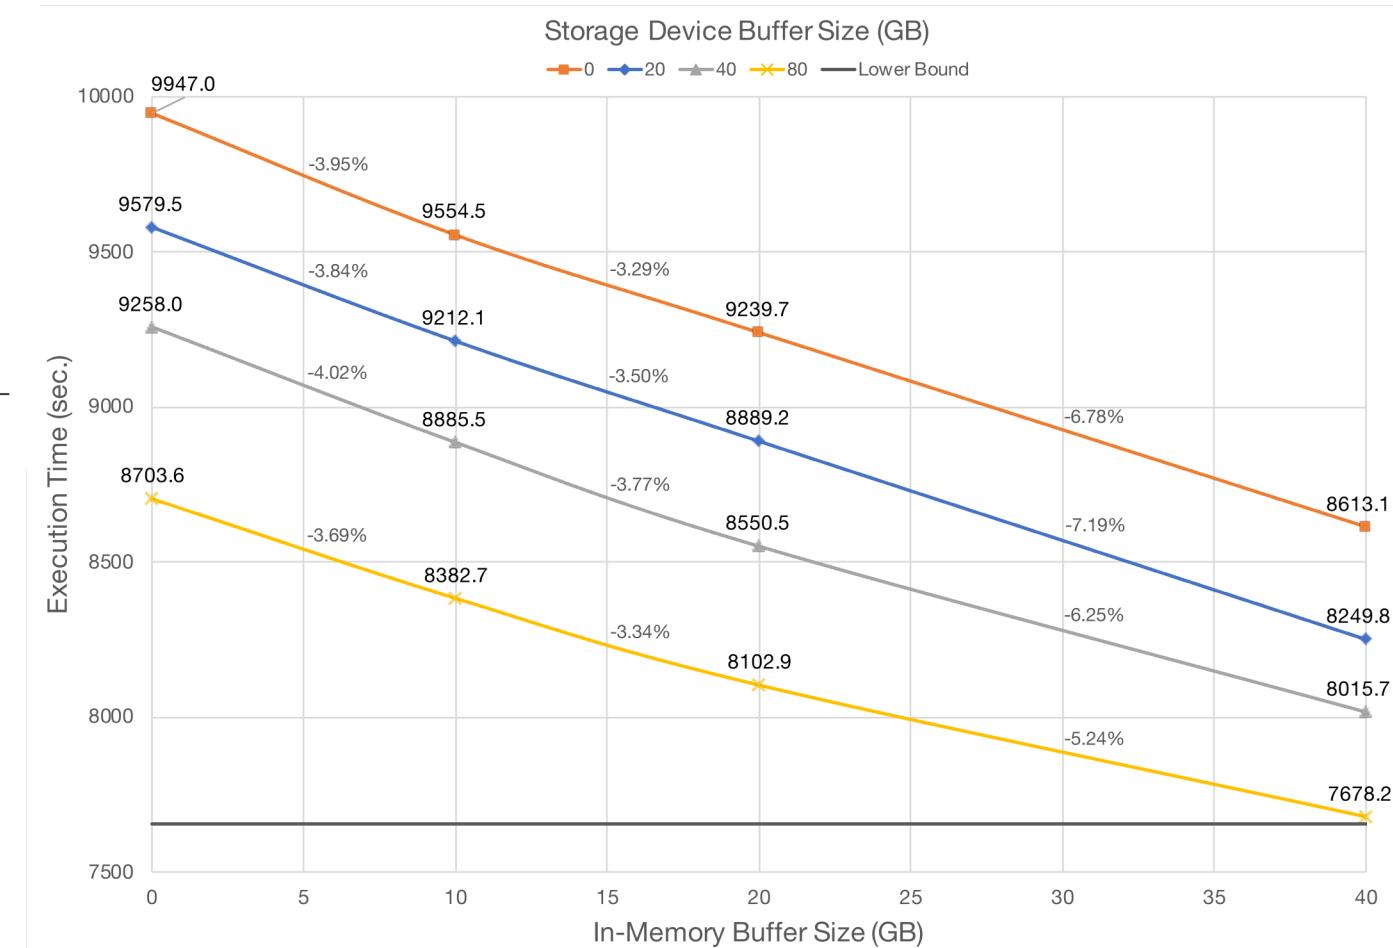
## 4. Larger datasets



# Performance Simulation: Design Validation & Environment Evaluation

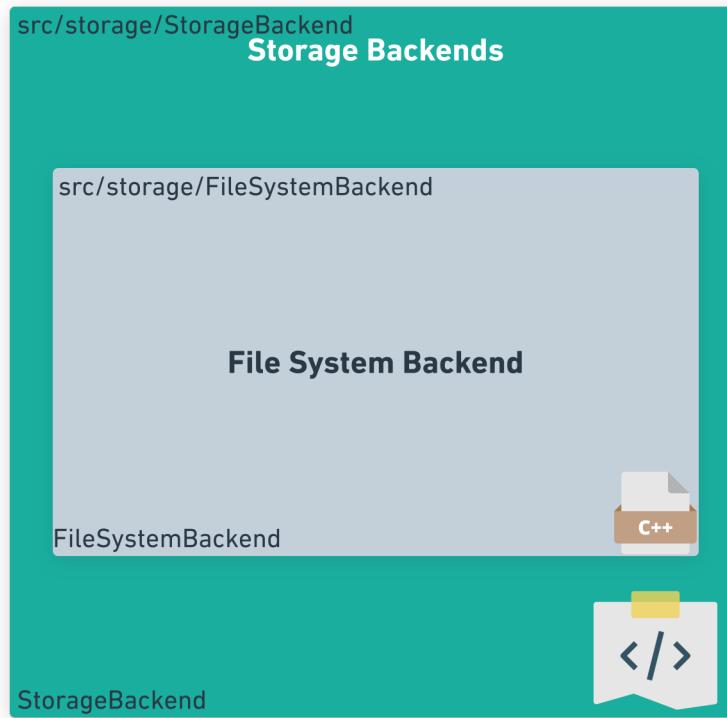
Policy	Runtime (s)
Lower Bound	3,825.66
<b>HDMLP</b>	<b>3,836.12</b>
LBANN In Memory Store (Dynamic)	4,051.12
DeepIO (Opportunistic)	4,432.90
DeepIO (Ordered)	4,433.27
LBANN In Memory Store (Preloading)	4,433.27
Locality-Aware Data Loading	4,433.27
Parallel Data Staging	4,972.65
StagingPoolPrefetcher	4,982.27
Naive	15,130.54

**Scenario 2 Results.**



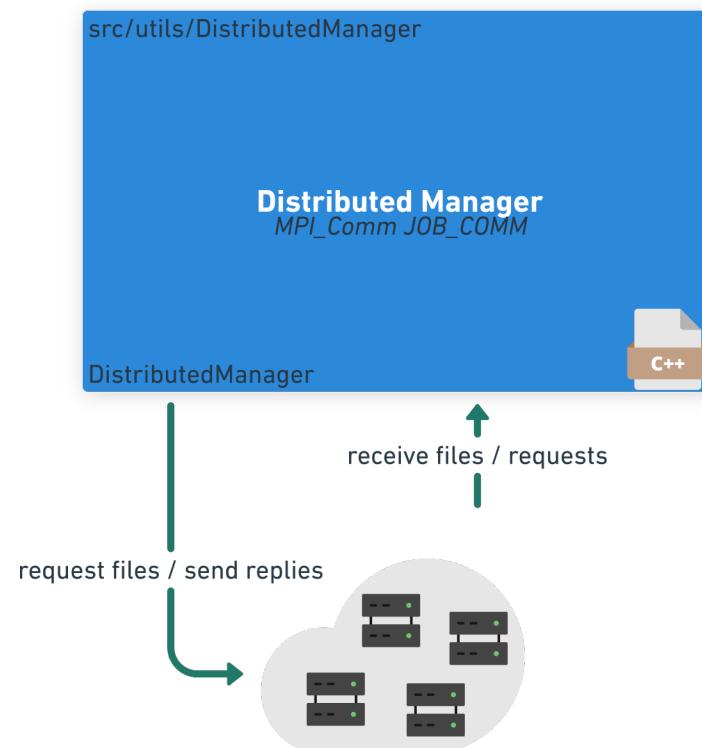
# System Implementation: Challenge 1 (Dataset Scalability)

- Arbitrary storage backends with minimal assumptions



# System Implementation: Challenge 2 (Node Scalability)

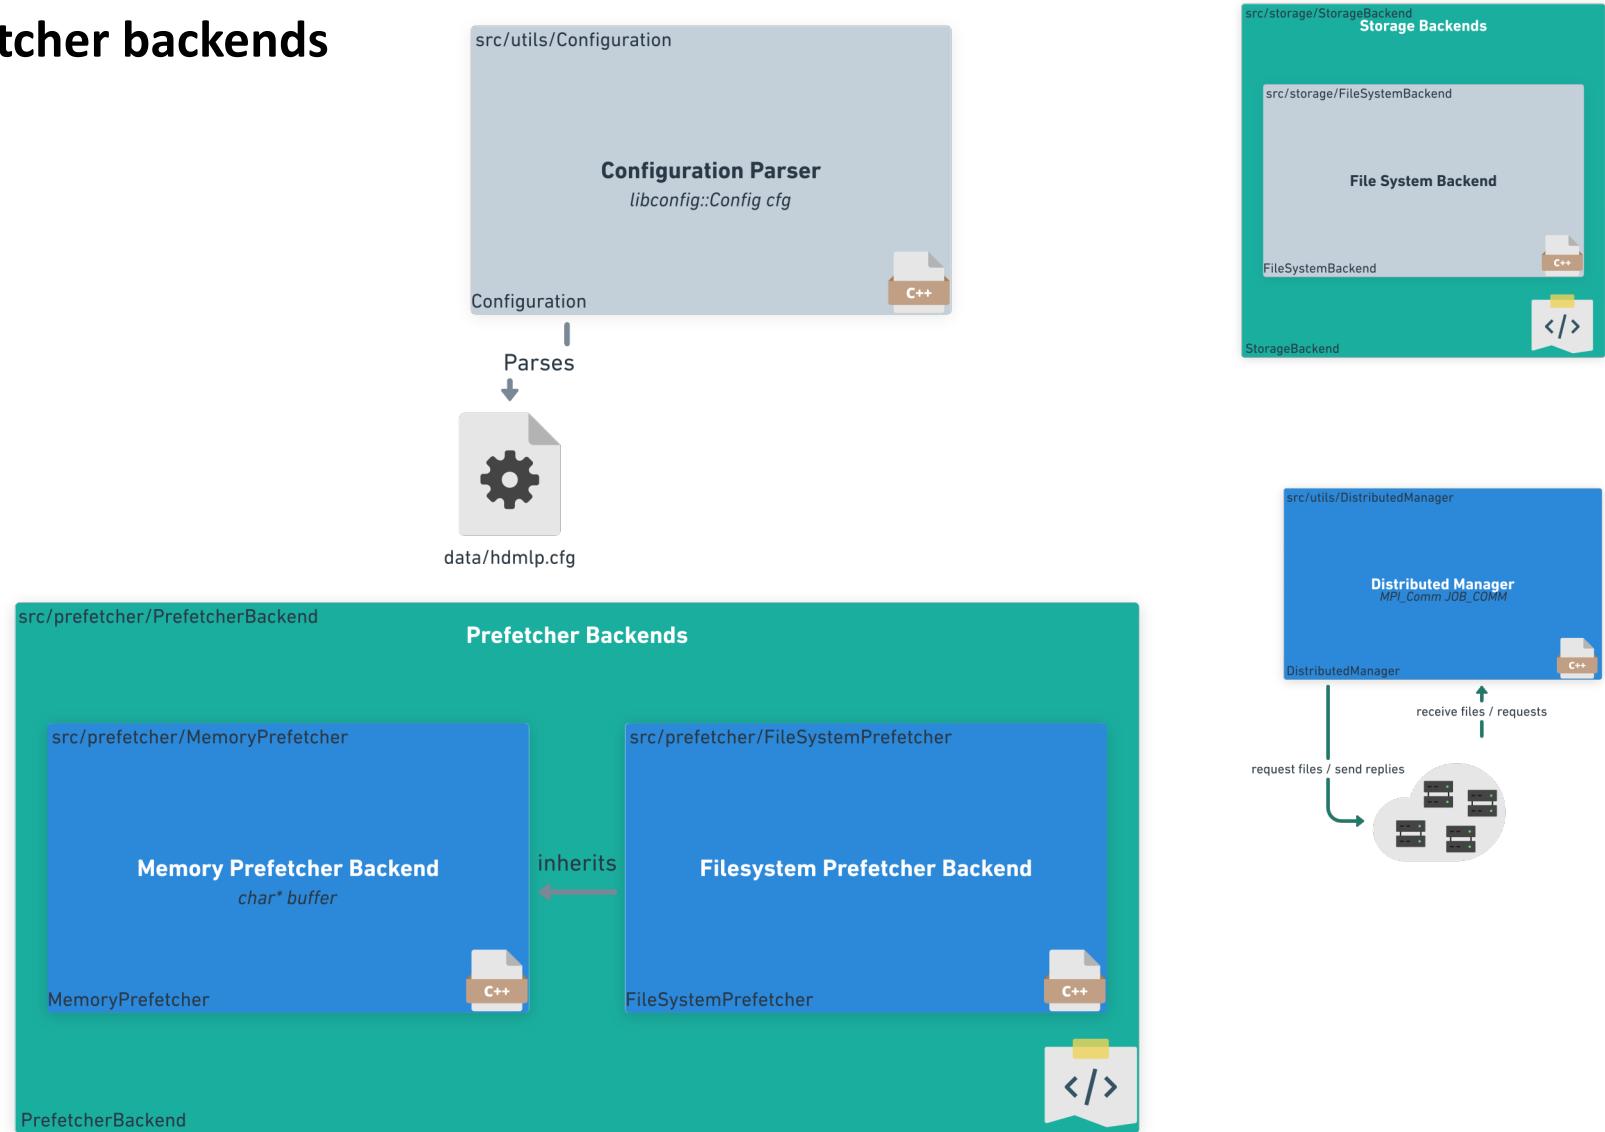
- Distributed caching without metadata messages



# System Implementation: Challenge 3 (Configuration Independence)

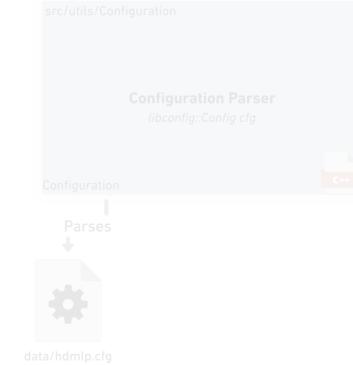
## Configurable and extensible prefetcher backends

```
1  b_c = 10000 # Network bandwidth for node communication [MB/s]
2  b_fs = 2000 # Network bandwidth for FS communication [MB/s]
3  distributed_threads = 2 # Number of threads for file distribution
4  checkpoint = true # Store label / size information in checkpoint file
5  checkpoint_path = "/tmp"
6
7  storage_classes: ( {
8      capacity: 1024 # [MB]
9      threads: 4 # Number of threads for prefetching
10     bandwidth: ( {
11         { threads = 1; bw = 16550; },
12         { threads = 2; bw = 21164; },
13         { threads = 3; bw = 21186; },
14         { threads = 4; bw = 21415; }
15     );
16 },
17 {
18     capacity: 2048 # [MB]
19     backend: "memory"
20     threads: 2 # Number of threads for prefetching
21     bandwidth: ( {
22         { threads = 1; bw = 16550; },
23         { threads = 2; bw = 21164; },
24         { threads = 3; bw = 21186; },
25         { threads = 4; bw = 21415; }
26     );
27 },
28 {
29     capacity: 10240 # [MB]
30     backend: "filesystem"
31     backend_options: {
32         path = "/tmp/hdmlp";
33     }
34     threads: 2 # Number of threads for prefetching
35     bandwidth: ( {
36         { threads = 1; bw = 16550; },
37         { threads = 2; bw = 21164; },
38         { threads = 3; bw = 21186; },
39         { threads = 4; bw = 21415; }
40     );
41 });
42 }
```



# System Implementation: Challenge 5 (Ease of Use)

## ■ Easy-to-use Python library



## ■ PyTorch datasets / dataloader



```
dataset = ImageFolder(data_dir, data_transforms) PyTorch
dsampler = DistributedSampler(dataset, num_replicas=n, rank=node_id)
dataloader = DataLoader(dataset, batch_size, sampler=dsampler)
```

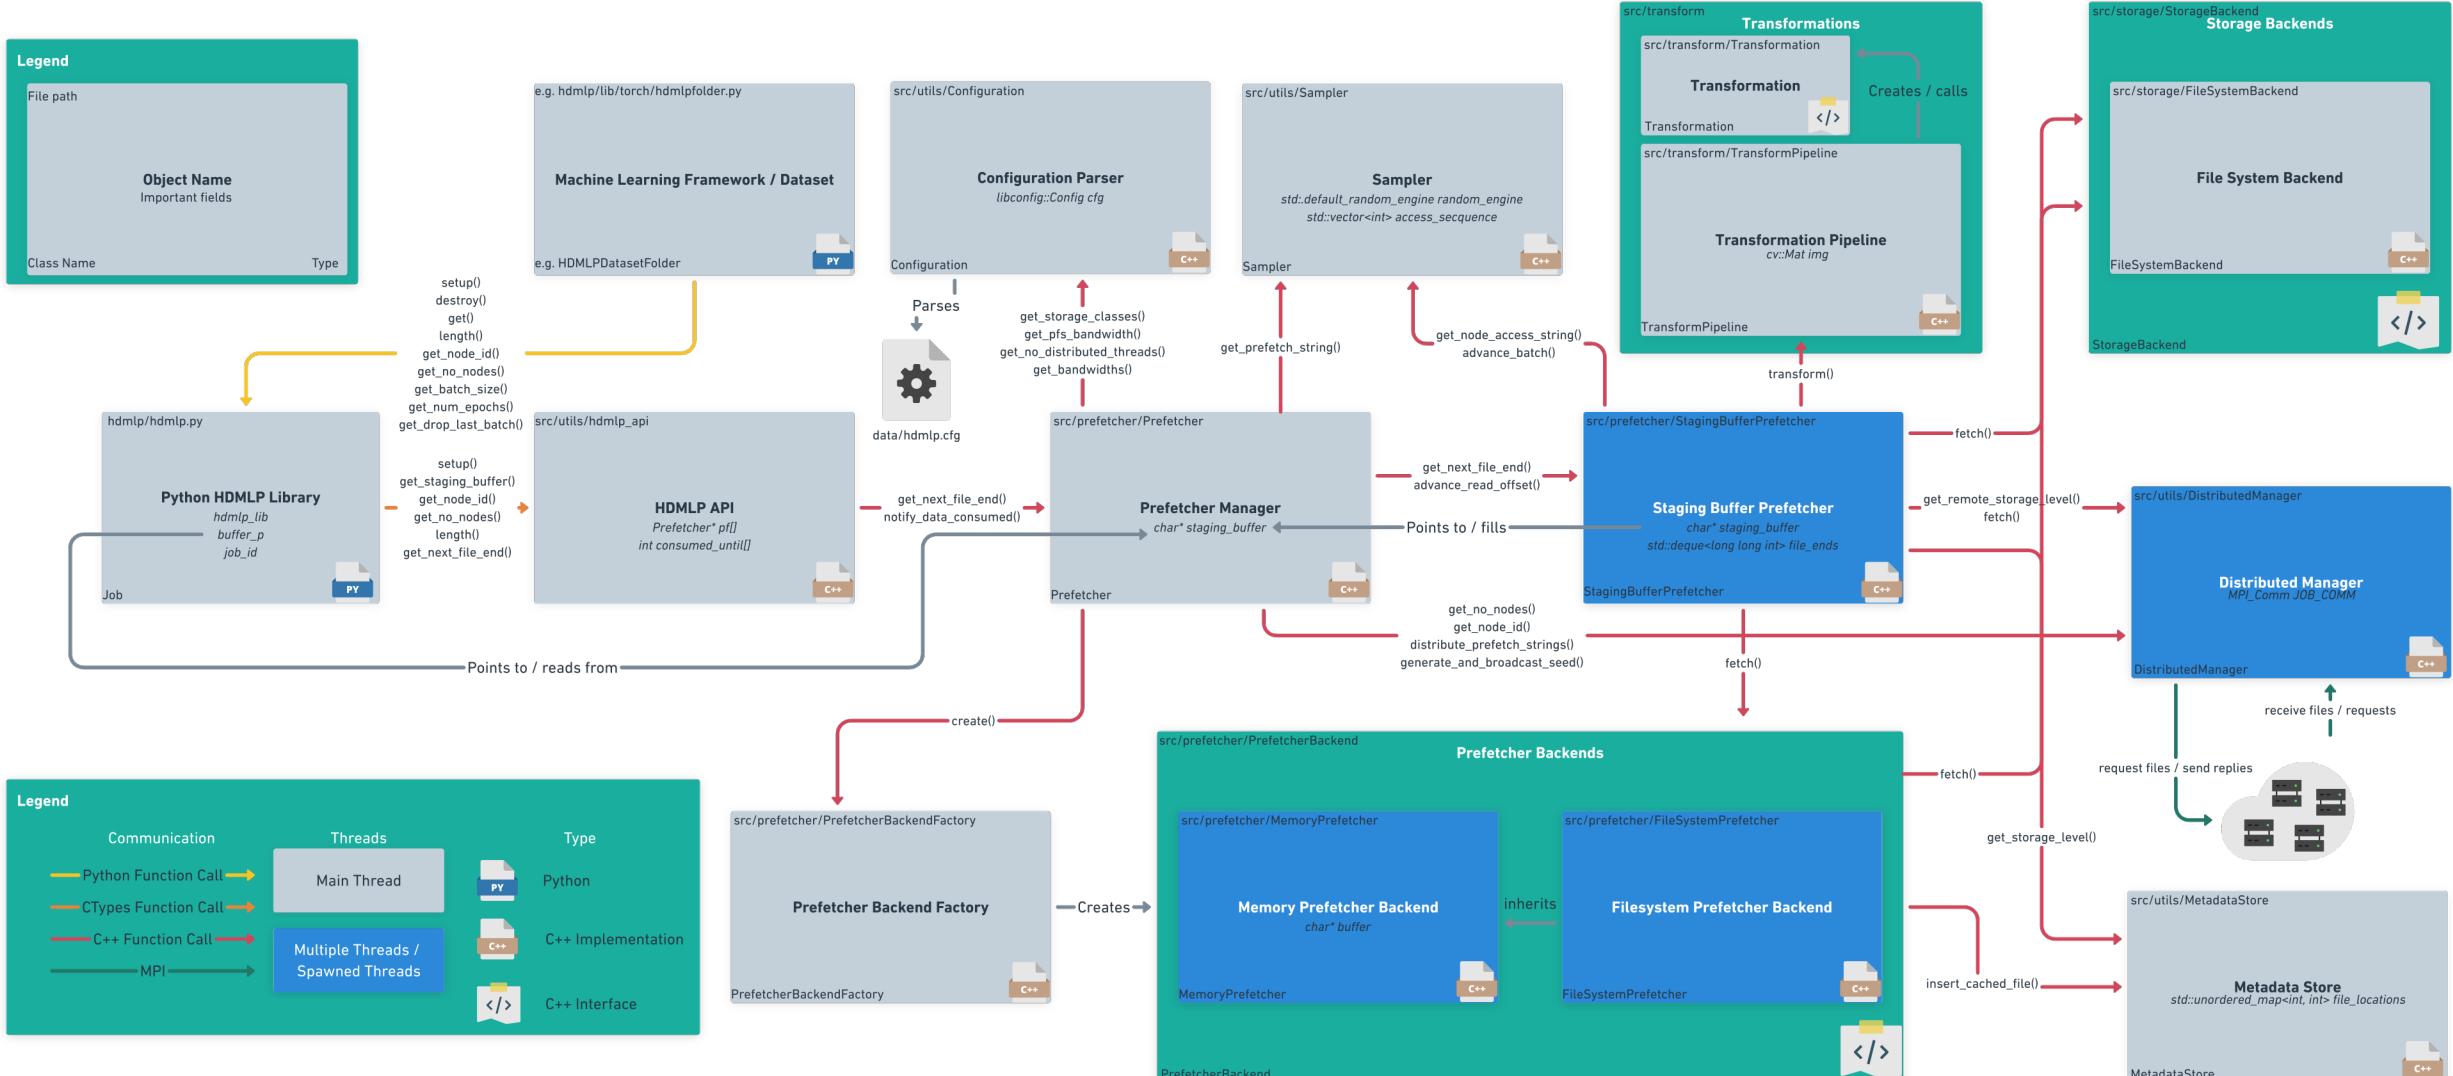
```
job = Job(data_dir, batch_size, num_epochs, 'uniform', drop_last_batch)
dataset = HDMLPImageFolder(data_dir, job, data_transforms)
dataloader = HDMLPDataLoader(dataset) HDMLP
```



# System Implementation: Zero-Copy Transformation Pipeline



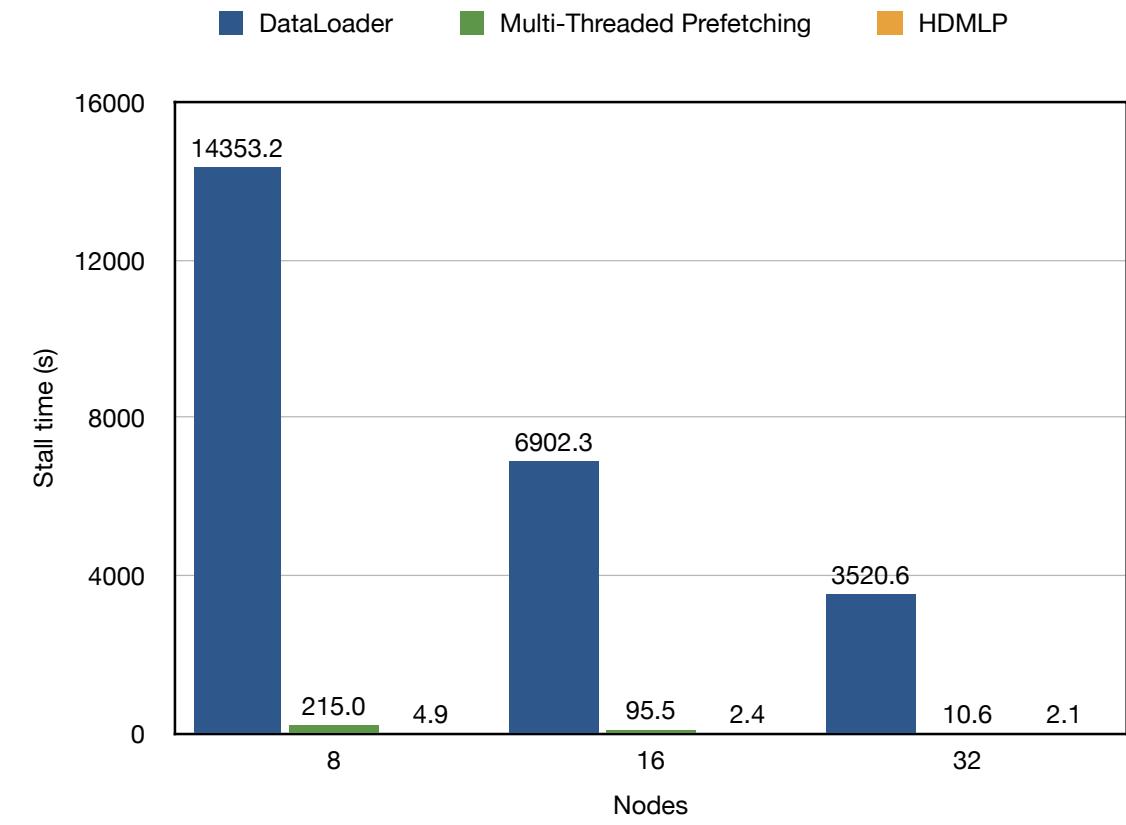
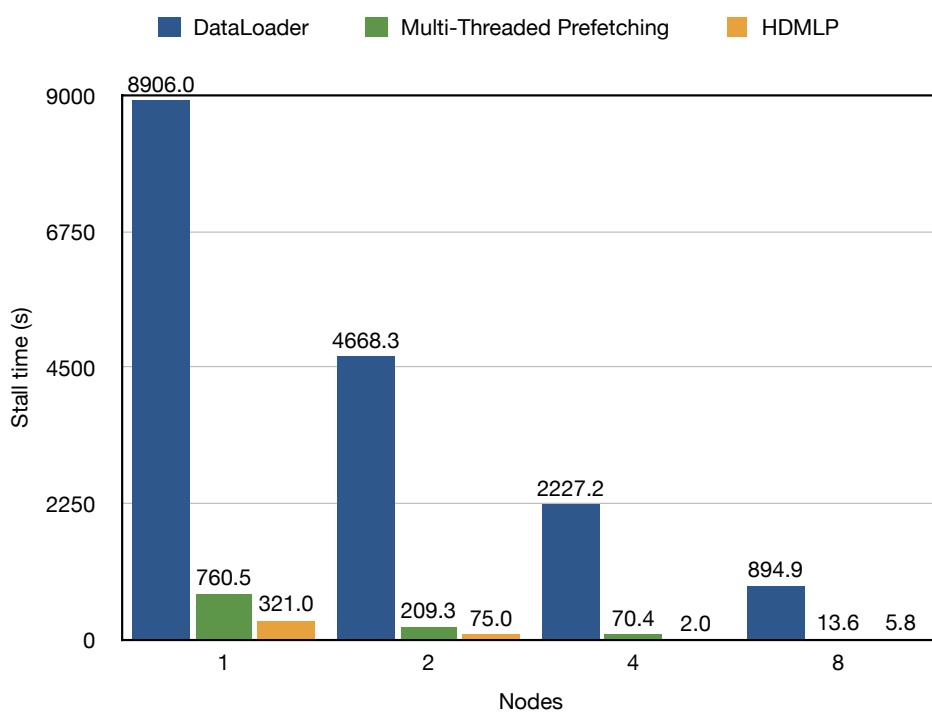
# System Implementation



# Evaluation

Nodes	Speedup	Strong Scaling Speedup	Strong Scaling Efficiency
1	2.433	1	1
2	2.408	2.005	1.002
4	4.057	7.192	1.798
8	2.147	7.640	0.955
16	2.298	17.987	1.124
32	2.433	40.146	1.255
64	2.476	85.591	1.337

Speedup, Strong Scaling Speedup, and Efficiency of HDMLP.



# Discussion, Next Steps

<https://github.com/spcl/hdmlp>



Approach	Dataset Scalability	Node Scalability	Configuration Independence	Model Accuracy	Ease of Use
tf.data with shuffle/prefetch [12]	✓	✗	✗	✗	✓
Double Buffering	✓	✗	✗	✓	✓
DeepIO [47]	✓	✓	✗	✗ <sup>a</sup>	✓
Parallel Data Staging [28]	✓	✓	✗	✗	✗
LBANN Distributed Data Store [21]	✗	✓	✗	✓	✗
Locality-Aware Data Loading [44]	✓	✓	✗	✓ <sup>b</sup>	? <sup>c</sup>
HDMLP	✓	✓	✓	✓	✓

- Provide loaders for other ML frameworks
- Extend HDMLP (storage backends, prefetcher backends, transformations, ...)